

DataEng S24: Data Validation Activity

High quality data is crucial for any data project. This week you'll gain experience with validating a real data set provided by the Oregon Department of Transportation.

Due: this Friday at 10pm PT

Submit: Make a copy of this document and use it to record your results. Store a PDF copy of the document in your git repository along with any needed code before submitting using the in-class activity submission form.

A. [MUST] Initial Discussion Question - Discuss the following question among your working group members at the beginning of the week and place your responses in this space. Or, if you have no such experience with invalid data then indicate this in the space below.

Have you ever worked with a set of data that included errors? Describe the situation, including how you discovered the errors and what you did about them.

Response:

There were anomalies in the data, but since the dataset was huge, discarding the abnormalities was not too much of a big deal.

The data set for this week is [a listing of all Oregon automobile crashes on the Mt. Hood Hwy \(Highway 26\) during 2019](#). This data is provided by the [Oregon Department of Transportation](#) and is part of a [larger data set](#) that is often utilized for studies of roads, traffic and safety.

Here is the available documentation for this data: [description of columns](#), [Oregon Crash Data Coding Manual](#)

Data validation is usually an iterative three-step process.

- A. Create assertions about the data
- B. Write code to evaluate your assertions.
- C. Run the code, analyze the results and resolve any validation errors

Repeat this ABC loop as many times as needed to fully validate your data.

B. [MUST] Create Assertions

Access the crash data, review the associated documentation of the data (ignore the data itself for now). Based on the documentation, create English language assertions for various properties of the data. No need to be exhaustive. Develop one or two assertions in each of the following categories during your first iteration through the ABC process.

1. *existence* assertions. Example: “Every crash occurred on a date”
 - Every crash has an associated county with it
2. *limit* assertions. Example: “Every crash occurred during year 2019”
 - Every crash occurred during a day in the week 1-7
3. *intra-record* assertions. Example: “If a crash record has a latitude coordinate then it should also have a longitude coordinate”
 - If a crash record has a latitude coordinate then it should also have a longitude coordinate
4. Create 2+ *inter-record check* assertions. Example: “Every vehicle listed in the crash data was part of a known crash”
 - Every crash that involves multiple vehicles should have the same respective dates for the crash.
 - Every crash has a crash serial number
5. Create 2+ *summary* assertions. Example: “There were thousands of crashes but not millions”
 - Most of the crashes involved multiple vehicles
 - Most of the crashes resulted in vehicular and property damages, not fatality.
6. Create 2+ *statistical distribution assertions*. Example: “crashes are evenly/uniformly distributed throughout the months of the year.”
 - Crashes are more frequent to weekdays rather than weekends
 - You can see peak crashes for rush hours (traffic hours)

These are just examples. You may use these examples, but you should also create new ones of your own.

C. [MUST] Validate the Assertions

1. Study the data in an editor or browser. Study it carefully, this data set is non-intuitive!.
2. Write python code to read in the test data. You are free to write your code any way you like, but we suggest that you use pandas’ methods for reading csv files into a pandas Dataframe.
3. Write python code to validate each of the assertions that you created in part A. The pandas package eases the task of creating data validation code.

4. If needed, update your assertions or create new assertions based on your analysis of the data.

D. [MUST] Run Your Code and Analyze the Results

In this space, list any assertion violations that you encountered:

- Assertion 4 failed
- Assertion 6 failed
- Assertion 7 failed
-
-

For each assertion violation, describe how to resolve the violation. Options might include:

- revise assumptions/assertions
- discard the violating row(s)
- Ignore
- add missing values
- Interpolate
- use defaults
- abandon the project because the data has too many problems and is unusable

No need to write code to resolve the violations at this point, you will do that in step E.

I changed the logic for the Assertion 4, instead of checking multiple values associated with date, I checked the serial number instead since they should all have the same one as well.

Assertions 6 and 7 failed because I undershot the amount of people dying and I thought that the crashes would involve more than one vehicle most of the time.

However in my notebook, I altered the code in place, so the previous assertions that did not work are not there.

So in summary, I changed the value for the assertion for number 4, checking only the serial number. For assertions 6 and 7 I ended up changing the assumption itself since I was wrong in my survey of the data.

E. [SHOULD] Learn and Iterate

The process of validating data usually gives us a better understanding of any data set. What have you learned about the data set that you did not know at the beginning of the current ABC iteration?

My assertions were wrong. I did not think that most of the crashes resulted in fatalities. I also thought that a crash meant two vehicles needed to be involved but that's a collision, a type of crash. You can crash into a pole and it would be a one vehicle crash.

Next, iterate through the process again by going back through steps B, C and D at least one more time.

F. [ASPIRE] Resolve the Violations and Transform the Data

For each assertion violation write python code to resolve the violation according to your entry in the "how to resolve" section above.

Output the validated/transformed data to new files. There is no need to keep the same, awkward, single file format for the data. Consider outputting three files containing information about (respectively) crashes, vehicles and participants.